

Cloud Recovery

Sprint 3

Carlton Duffett, Deema Kutbi, Emilio Teran
Konstantino Sparakis, Minhan Xiang

ReClo Overview

Our Purpose:

Provide an affordable, simple means for recovery of failed servers

Our Clients:

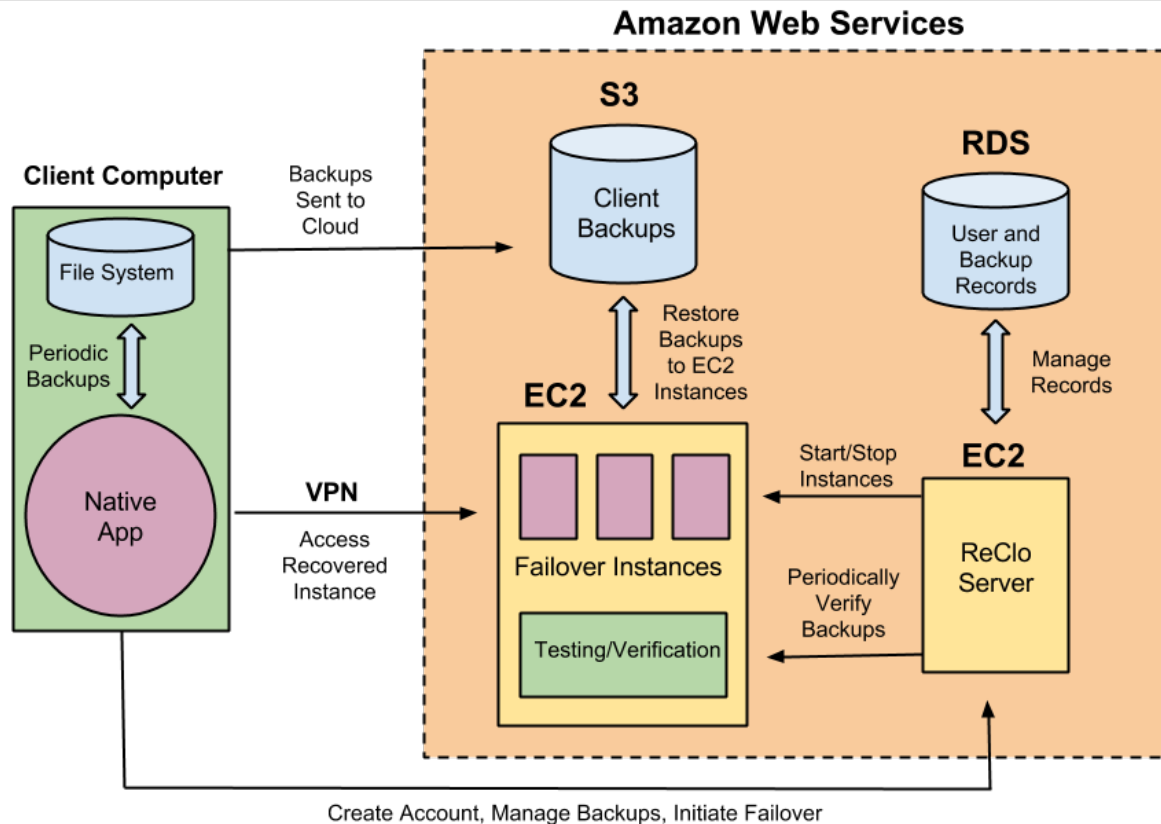
Small to medium-sized businesses, like franchise stores or doctors offices

Key Features:

Backup Manager performs regular backups and pushes them to the cloud

Recovery Manager starts a new EC2 instance from a backup file and establishes a VPN connection, allowing users to access their data

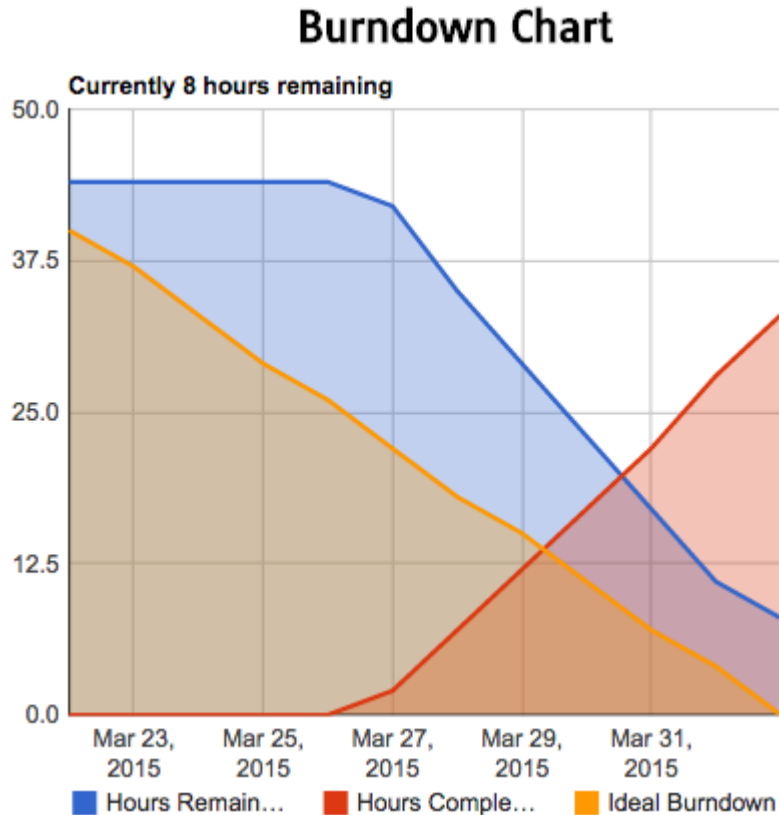
Architecture



This Sprint: Our Goals

- Functional Backup Manager and Recovery Manager UIs
- Incremental Backups
- Complete API (node.js) and client library (C#)
- Upload backups to S3
- Better estimation of tasks and better burndown pace

This Sprint: Burndown



Our best burndown yet.

Took a while to do research and implement features before anything was completed.

Getting close to the finish!

This Sprint: Trello

The screenshot shows a Trello board titled "EC500 - Sprint 3". The board is set to "Private" and has a "Burndown Chart" link. It is divided into three main columns: "Sprint Backlog", "In Progress", and "Done".

- Sprint Backlog:** Contains a placeholder "Add a card...".
- In Progress (11 cards):**
 - Card 1: "C# upload VHD back up file to S3" with 3 points, assigned to KS.
 - Card 2: "Synchronizing Backup Manager with C# client library and Powershell backup calls" with 8 points, assigned to DK, ET, and MX.
- Done (28 33 cards):**
 - Card 1: "Finish C# client library functions" with 4/4 points, 5 points, 8 points, assigned to KS.
 - Card 2: "Update website content" with 4/4 points, 2 points, 2 points, assigned to a user.
 - Card 3: "Recovery Manager GUI" with 3 points, 5 points, assigned to DK and ET.
 - Card 4: "Documentation" with 2/2 points, 2 points, 2 points, assigned to KS.
 - Card 5: "Finish Node.js server API calls" with 5/5 points, 8 points, 8 points, assigned to a user.
 - Card 6: "Incremental Backups" (no points or assignees visible).

What We Accomplished:

- **Fully-Tested API** and client for user_auth, backups
- **Expanded API** for starting, stopping EC2 instances
(still working on starting instances from backups)
- **Incremental Backups** from Backup Manager
(still working on uploading these backups to S3)
- Recovery Manager UI
- Cron jobs to periodically maintain our resources

Recovery Manager Application

ReClo

Recover

Recovery source:

File version to recover: ☐ Latest ☐ Specific Date:

Recovery destination:

What to recover:

Description:

ReClo

Recover

Recovery Source:

File Version to Recover: ☒ Latest ☐ Specific Date:

Recovery Destination:

What to Recover:

Description:

Backup Manager

Incremental Backups

Tool: Cobian Backup 8

- Supports incremental backups
- Supports command line operations
- Free
- Backed-up files are kept in original format

Incremental Backups

What we do:

- Initialize Cobian backup task
- Execute Cobian task from our backup manager
- Convert the backed-up file into .vhd format

RESTful API

Added API calls to start, stop,
and get information on EC2
instances



Amazon EC2

Uses node.js (lightweight,
asynchronous JavaScript) for the server.

C# client library for use with the applications.

Cron Jobs

ManageSessions()

Periodically scans our database and deactivates expired session tokens

FreeInstances()

Permanently terminates any stopped instances that have been idle for more than 2 weeks

Forever.js

Easy to Use:

Perfect tool for running a stand alone web server. No need for Apache or nginx.

Fault tolerant:

Keeps our server running continuously and automatically restarts the server when it exits unexpectedly.

API Documentation

Response Error Codes

List of possible error codes that the API will return (user_auth, backups, recovery). All error codes

system-wide:

101 cannot connect to database
102 invalid token

/register:

201 user with that email already exists
202 invalid username format (must be > 4 or
203 invalid email format
204 invalid password format (at least 6 char)
205 failed to create S3 bucket when new user

/login:

206 user not found
207 password does not match
208 invalid email format
209 invalid password format

/logout:

211 failed to deactivate token

/backups/getBackupList():

301 failed to get list of backups
302 no backups exist for user

/backups/startUpload():

303 unable to get temporary S3 credentials

/backups/completeUpload():

304 failed to complete upload
305 no upload found
306 failed to get backup information
307 could not verify backup

/backups/deleteBackup():

308 failed to obtain file_name to delete
309 no backup to delete
310 failed to delete backup

startUpload()

Obtains temporary (12 hour) AWS S3 cred

Format of Call:

POST /backups/uploads/:user_id?token=

Request Parameters:

1. user_id passed as part of the
2. token passed as part of the
3. file_name passed as part of the
4. file_size passed as part of the

Example Request:

52.11.1.237:3000/backups/uploads/c-413ed6c7938c47889b43d32f2c525aae

Response Content:

1. upload_id unique identifier for the backup
2. credentials object
 - a. AccessKeyId public
 - b. SecretAccessKey secret
 - c. SessionToken temp each /
 - d. Expiration localiz

Example Response:

```
{
  "message": "Obtained temporary credentials",
  "upload_id": "9a67bf35793c477e",
  "credentials": {
    "AccessKeyId": "ASIAIN33SQ",
    "SecretAccessKey": "0/hcsj9",
    "SessionToken": "AQoDYXZlZDID////////wEagAN3aGJwOYa7yWcRLgPHMzxfCZGTvNcwZ09nZl0S3jSncrhBOWCrG913VHMEhtwFY5=f+5AYCqIH6WicJHrF17F/5sF6DCkuQtIKSlnw24GdazBX5Id8qAKAJxTxmw0TZqo9UGLVGFhy/n1sln/L33bvy51168tvy4dMUNOKf1jHBKH1180h2v3TtdIav9eDRIPAp/9NSXbuHRRKPSY-XURNN1j9ckh8gpp5u94AU=",
    "Expiration": "2015-03-23T"
  }
}
```

Backups API Calls

getBackupList()

Get a list of all backups current stored for a given user

Format of call:

GET /backups/:user_id?token=

Request Parameters:

1. user_id passed as part of the URL
2. token passed as part of the URL

Example Request:

52.11.1.237:3000/backups/c2c7281d-661c-token=413ed6c7938c47889b43d32f2c525aae

Response Content:

backups array of backup objects as follows:
1. backup_id unique identifier of each backup
2. file_size file size in MB
3. file_name file name
4. date_created UTC timestamp as a string

Example Response:

```
{
  "message": "Backups obtained successfully",
  "backups": [
    {
      "backup_id": "073b691c7f01",
      "file_size": 28.3,
      "file_name": "backup01.vhd",
      "date_created": "Thu, 19 Mar 2015 14:00:00 GMT"
    },
    ...
  ]
}
```

Node.js API Documentation

Overview

All responses are in JSON

If an error occurs, the server responds with status(500) and returns: {"error": error_code, "message": "error message..." }

A list of potential error codes is provided in the **Error Codes** section.

If the operation is successful, the server responds with status(200) and returns: {"message": "success_msg here", ... }

The "message" is followed by any operation-specific content.

Usage

Most API calls require a valid session token as part of the URL. Tokens are 32-character alphanumeric strings. Tokens are always passed in the URI query, for example:

GET ip_addr:port/api_call?token=<your token here>

Session tokens are invalidated after **60 days**. We assume that within 60 days our users will be able to restore their failed hardware. Users or course may renew their session.

If a user stops an instance, that instance is held for a period of **2 weeks**. If an instance remains stopped for more than 2 weeks, it is permanently terminated to free space.

API Client Library

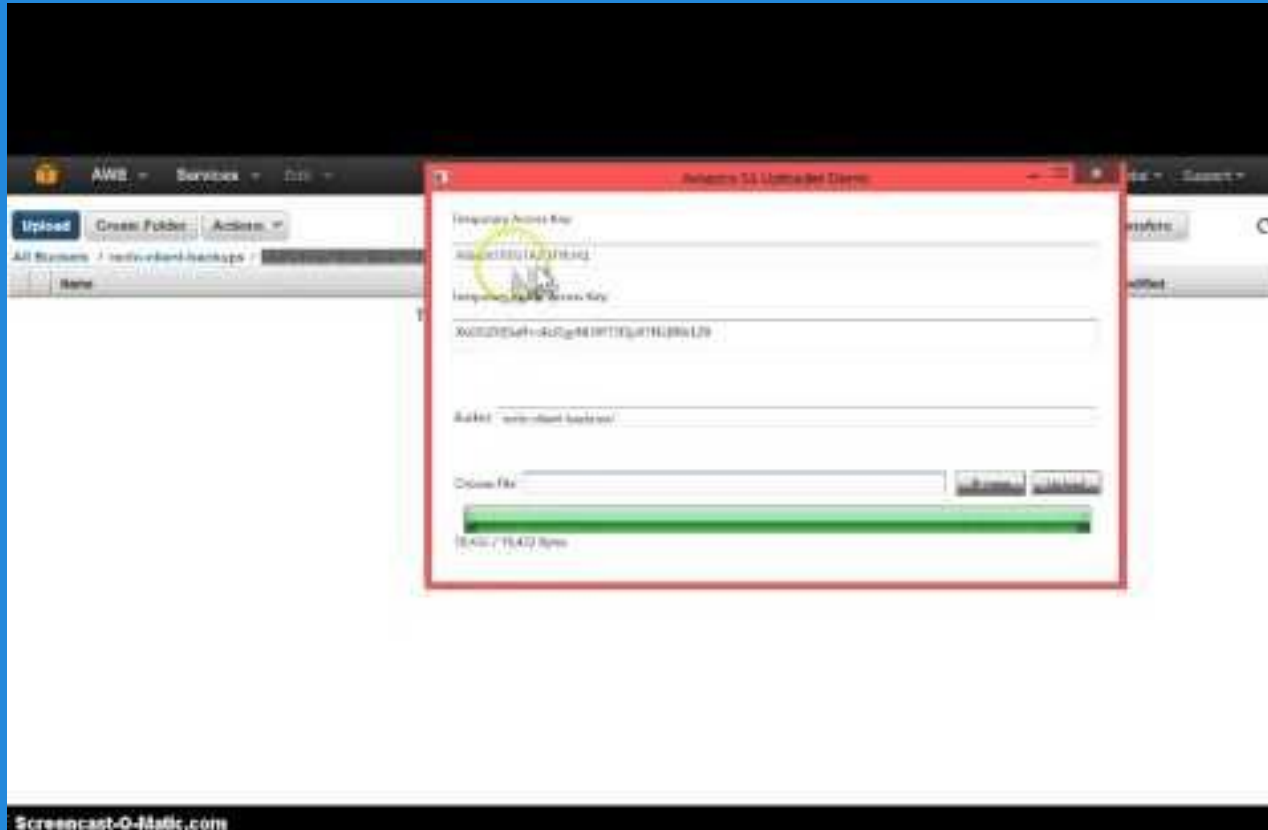
- `HttpMethods.cs` and `RecloApiCaller.cs`
- All calls to API are asynchronous
- Finished functions for all backup API calls
- Tested all functions except for DELETE
- `mc.sh`
- Simplified and documented

S3 Upload

- Temporary Access token to user
- Upload is done from the C# client using AWS SDK
- Upload is asynchronous
- Needs to be merged with GUI and api client code



S3 Upload Video Demo



A Perspective

Language	Lines of Code Written
Node.js	1610
C#	779
HTML & CSS	220
TOTAL	2609

2609 lines of ...

1. Debugged
2. Running
3. Tested

... code written!

Next Sprint: Goals

- Upload backups to S3
- Start new EC2 instances from vhd backups
- VPN Connection to restored instances
- Finish Backup Manager back-end
- Recovery Manager back-end